




भारतीय संगणक एवं प्रौद्योगिकी संस्थान

INDIAN INSTITUTE OF COMPUTING AND TECHNOLOGY

AFFILIATED: I-STEM, OFFICE OF THE PRINCIPAL SCIENTIFIC ADVISER TO THE GOVERNMENT OF INDIA

AI-Driven Phishing Email Detection Using NLP
(Natural Language Processing)



Project Report
15th June to 30th July 2026

Student Name Anand krishna G R Nair

Institute Name: Jecrc University

Institute Roll no.: 23BCON1613

Enrollment no.: 596563

AI-Driven Phishing Email Detection Using NLP

Anand Krishna G R Nair

Department of Computer Science and Engineering
Indian Institute of Computing and Technology (IICT), New Delhi — 110092

Registration No: 596563

Email: anandkrishnag.rnair@gmail.com

ORCID: 0009-0008-0757-1566

Manuscript received 1 July 2026; revised 10 July 2026; accepted 25 July 2026.

Abstract — Modern electronic mail infrastructure remains highly vulnerable to adversarial phishing attacks that bypass technical gateways by exploiting human cognitive trust. This project implements a high-accuracy, low-latency machine learning system for real-time phishing classification. Using a benchmark dataset of 17,535 preprocessed email samples, we construct a 5,112-dimensional hybrid feature matrix integrating sublinear TF-IDF n -grams (5,000 features), custom dense Word2Vec CBOW embeddings (100 features), and 12 structural metadata metrics (such as URL counts, urgency keyword matches, and capitalisation ratios). Five candidate classifiers—Multinomial Naive Bayes, Logistic Regression, Random Forest, Multilayer Perceptron (MLP), and a Soft-Voting Ensemble—were evaluated across 5-fold Stratified Cross-Validation. On an independent 20% holdout test set (3,507 samples), the Soft-Voting Ensemble achieved an Accuracy of 98.52%, a Precision of 98.09%, a Recall of 97.94%, and an F1-Score of 98.02%. Benchmarking demonstrates an average inference latency of 2.10 ms per email, validating CPU gateway deployability without expensive GPU infrastructure.

Index Terms — Cybersecurity, Email Classification, Feature Engineering, Machine Learning, Natural Language Processing, Phishing Detection, Soft-Voting Ensemble, TF-IDF, Word2Vec.

INTRODUCTION & THREAT LANDSCAPE

Threat Evolution and Enterprise Risk

Electronic mail serves as the foundational communication protocol for business, academic, and governmental enterprises worldwide. Despite decades of security innovations, cyber threat intelligence reports confirm that over 90% of organizational security incidents originate from malicious email messages. Unlike network-level exploits that target unpatched software vulnerabilities, phishing attacks exploit human cognitive biases—such as perceived urgency, fear of penalty, authority impersonation, and misplaced trust.

Early phishing prevention relied primarily on static signature blacklists and manual rule-based spam filters. While static blacklists are efficient at blocking known malicious sender addresses and domain URLs, they are inherently incapable of identifying zero-day phishing campaigns, rapid domain rotation, or spear-phishing messages tailored to specific executives. As attackers adopt sophisticated evasion techniques—including text obfuscation inside HTML comment tags, homograph domain spoofing, and clean vocabulary styling—conventional rule-based security gateways suffer unacceptably high false-negative rates.

Modern enterprise security environments process tens of millions of inbound email messages daily. Security operations centers (SOCs) face unprecedented challenges in maintaining real-time filtering without degrading mail server delivery performance. A delayed message can disrupt time-sensitive financial transactions or emergency corporate communications, while a false negative can expose an entire internal network to ransomware, credential theft, and financial fraud.

Evasion Tactics and Technical Challenges

Modern attackers intentionally structure messages to confuse standard bag-of-words classifiers. Common evasion tactics include:

- 1) *HTML Payload Obfuscation*: Injecting invisible HTML comment tags (e.g., `U<!-- comment -->RGENT`) or transparent text colors between letters of urgent words so human readers see 'URGENT' while basic text scrapers see disconnected characters.
- 2) *Hyperlink Redirection*: Embedding clean visible hyperlink anchor text (e.g., 'http://bank.com') while routing the underlying HTML href attribute to external malicious IP addresses or compromised short-link services.
- 3) *Semantic Urgency Mimicry*: Crafting polished, grammatically flawless text that mimics legitimate corporate human resources, legal, or IT support notifications, eliminating the spelling errors historically relied upon by

naive filters.

- 4) *Rapid Domain Rotation*: Registering short-lived domains that remain active for only hours, rendering static threat intelligence blacklists obsolete before threat feeds update.

Research Objectives and Systematic Scope

This project addresses these technical challenges by designing, implementing, and empirically evaluating an end-to-end machine learning pipeline capable of real-time phishing classification. To achieve scalable defense, five core research objectives guide this investigation:

- **RQ1 (Feature Fusion Efficacy)**: Does integrating 5,000 sublinear TF-IDF n-grams with 100 dense Word2Vec CBOW embeddings and 12 structural metadata features outperform single-source feature spaces?
- **RQ2 (Ensemble Synergies)**: How effectively does a Soft-Voting Ensemble reconcile individual model decision boundaries across linear (Logistic Regression), tree-based (Random Forest), neural (MLP), and probabilistic (Naive Bayes) paradigms?
- **RQ3 (Gateway Latency Constraints)**: Can multi-modal feature extraction and ensemble classification execute within a sub-5ms SLA on standard CPU hardware to support high-throughput mail server deployment?
- **RQ4 (Evasion Resilience)**: Do structural metadata indicators maintain classification confidence when text payloads undergo zero-day adversarial obfuscation?
- **RQ5 (Model Interpretability)**: What specific syntactic n-grams and structural formatting ratios drive model predictions across safe and malicious classes?

LITERATURE REVIEW & TAXONOMY

Historical Heuristics & Statistical Baselines

Automated email filtering has been researched extensively since the late 1990s. Sahami et al. (1998) pioneered Naive Bayes spam classification, demonstrating that word frequency distributions could separate commercial spam from legitimate correspondence. However, Naive Bayes models assume conditional feature independence given the class label—an assumption that fails when analyzing complex phishing phrases where multi-word combinations carry specific malicious intent.

Egozi and Verma (2018) extended lexical feature engineering by demonstrating that grammatical and structural formatting flags significantly improve classifier resilience against text obfuscation. Their work established that combining lexical metrics with structural formatting flags allows machine learning models to detect suspicious

message characteristics even when attackers deliberately alter their vocabulary.

Subsequent investigations by Zhang et al. explored support vector machines (SVMs) using non-linear radial basis function (RBF) kernels. While SVMs established strong marginal hyperplanes, their training computational complexity scales quadratically with dataset size, rendering them impractical for daily enterprise retraining across hundreds of thousands of samples.

Multi-Modal & Metadata Formulations

Subsequent research focused on multi-modal feature representations. Al-Subaee and Al-Zahrani (2020) proposed combining lexical TF-IDF body features with domain age, URL counts, and header metadata. Their empirical results showed that multi-source feature fusion drastically reduces false-positive rates compared to pure text models. Similarly, Gedam and Shende (2023) confirmed that capitalisation ratios, punctuation density, and urgency vocabulary scores serve as vital indicators for phishing classification.

Multi-modal feature spaces leverage the intrinsic structural syntax of email protocols. By analyzing character distributions, capitalisation frequency, and embedded link structures alongside natural language tokens, machine learning systems construct resilient multi-layered decision boundaries that withstand single-vector evasion tactics.

Deep Learning vs. Classical Ensembles

Recent research applied transformer language models (BERT, RoBERTa, DistilBERT) to email security (Salloum et al., 2021). While contextual transformers achieve strong benchmark results, their operational forward-pass latency (50–200 ms per email) and substantial GPU memory demands create severe bottlenecks in enterprise security gateways processing millions of emails daily. Bountakas and Xenakis (2021) also noted that deep models remain susceptible to subtle character-level perturbations. Consequently, hybrid feature engineering combined with classical ensembles (Random Forest, MLP, Soft Voting) represents the most practical, high-throughput solution for real-time gateway defense.

TABLE I. TAXONOMY COMPARISON OF PHISHING DETECTION METHODOLOGIES

Study / Ref.	Technique	Advantage	Limitation
Sahami (1998)	Naive Bayes + BoW	Fast, low memory	Assumes independence
Egozi (2018)	Lexical + Struct.	Obfuscation resilient	Manual feature design
Al-Subaee (2020)	TF-IDF + Metadata	Low false positives	Needs network telemetry
Salloum (2021)	BERT / Transformers	Deep context	High latency (50-200ms)
Our Work (2026)	TF-IDF + W2V + Meta	98.52% Acc, 2.1ms	Needs W2V retraining

SYSTEM ARCHITECTURE & METHODOLOGY

Modular Pipeline Topology

The system architecture follows a modular feed-forward pipeline engineered for high-throughput execution. As shown in Figure 1, the pipeline accepts raw email messages containing unstructured text and embedded HTML formatting, executes sequential sanitization, extracts multi-modal feature vectors, and evaluates class probabilities via a trained Soft-Voting Ensemble.

End-to-End Hybrid Phishing Email Detection Pipeline Architecture




FIGURE 1. End-to-end hybrid phishing detection pipeline architecture.

Data Ingestion & Sanitization Subsystem

Raw email files contain complex HTML markup, CSS, and script blocks. The sanitization module executes four sequential steps: BeautifulSoup parses raw strings, stripping script tags, style blocks, and comments while preserving link presence; text is normalized to lowercase, multiple whitespaces collapse, non-standard ASCII characters are removed, NLTK English stopwords are filtered, and duplicate email entries are dropped (leaving 17,535 clean samples) to prevent data leakage.

Hybrid Feature Matrix Construction

We construct a combined 5,112-dimensional feature vector derived from three complementary feature sources: a 5,000-dimensional sublinear TF-IDF n-gram vector space (unigrams, bigrams, trigrams), a 100-dimensional dense custom Word2Vec CBOW embedding vector derived via element-wise mean vector averaging, and a 12-dimensional structural metadata vector capturing formatting traits, urgency flags, and structural characteristics.

TABLE II. COMPOSITION OF 12 STRUCTURAL METADATA FEATURES

Feature Name	Type	Behavioral Indicator
url_count	Int	Count of HTTP/HTTPS URLs in body.
has_link	Binary	Flag indicating presence of any link.
email_address_count	Int	Count of embedded contact emails.
urgency_score	Int	Matches against 20 urgency words.
char_length	Int	Total character length of email.
word_count	Int	Total word count of token stream.
avg_word_length	Float	Average length of constituent words.
exclamation_count	Int	Count of exclamation marks (!).
question_count	Int	Count of question marks (?).
punctuation_density	Float	Ratio of punctuation to length.
caps_ratio	Float	Ratio of uppercase letters to length.
digit_ratio	Float	Ratio of numeric digits to length.

MATHEMATICAL FORMULATIONS & CLASSIFIER TOPOLOGY

Feature Space Derivations

Let $D = \{d_1, d_2, \dots, d_N\}$ represent the corpus of $N = 17,535$ cleaned email documents. For a term t in document d , sublinear TF-IDF scaling is computed as:

$$TF\text{-}IDF(t, d, D) = (1 + \log(tf(t, d))) \times \log(1 + N / df(t))$$

Sublinear term frequency scaling compresses repetitive word counts, preventing high-frequency terms from dominating classifications.

For Word2Vec dense vectors, each word w is mapped to a 100-dimensional vector $v_w \in \mathbb{R}^{100}$. The document vector v_d is obtained by computing the element-wise mean:

$$v_d = (1 / |d|) \times \sum_{w \in d} v_w$$

Dense Word2Vec vectors v_d and structural metadata vectors $m_d \in \mathbb{R}^{12}$ are normalized to $[0, 1]$ via MinMaxScaler and concatenated with sparse TF-IDF vectors $x_{tfidf} \in \mathbb{R}^{5000}$:

$$X_{final} = [X_{tfidf} \parallel \text{Scale}(V_{w2v}) \parallel \text{Scale}(M_{meta})] \in \mathbb{R}^{N \times 5112}$$

Classifier Formulations & Optimization

1) *Multinomial Naive Bayes*: Computes posterior probability with Laplace smoothing:

$$P(y = c | x) \propto P(y = c) \times \prod_{i=1}^K (x_i + \alpha) / (\sum x_i + \alpha K)$$

where $\alpha = 0.1$ is additive Laplace smoothing, and $K = 5112$ is the feature dimension.

2) *Logistic Regression*: Models class probability using the logit function:

$$P(y = 1 | x) = 1 / (1 + \exp(-(w^T x + b)))$$

Optimized using L2 regularization penalty $L(w) = \|w\|_2^2$ with parameter $C = 1.0$.

3) *Random Forest*: Ensemble of $T = 100$ unpruned trees evaluated via Gini Impurity:

$$\text{Gini}(p) = 1 - \sum_{c=0}^1 p_c^2$$

4) *MLP Neural Network*: Feed-forward network with one hidden layer of 100 neurons:

$$h = \text{ReLU}(W_1 x + b_1), y = \text{Sigmoid}(W_2 h + b_2)$$

5) *Soft-Voting Ensemble*: Computes decision boundaries by averaging predicted probabilities:

$$P_{\text{ensemble}}(y = c | x) = (1/M) \times \sum_{m=1}^M P_m(y = c | x)$$

EXPERIMENTAL RESULTS & EVALUATION

Dataset Demographics and Holdout Breakdown

The empirical benchmark corpus comprises 17,535 unique, preprocessed email samples. Class distribution breakdown reveals 10,977 legitimate (safe) emails, representing 62.6% of the total dataset, and 6,558 phishing emails, representing 37.4%. This natural baseline ratio provides a realistic representation of enterprise network traffic where legitimate communications predominate, while maintaining sufficient phishing instances to train non-trivial decision boundaries without requiring synthetic oversampling techniques (such as SMOTE).




FIGURE 2. Benchmark dataset class distribution (17,535 samples).

Stratified Cross-Validation and Metric Definitions

All candidate classifiers underwent 5-fold Stratified Cross-Validation on the 14,028 training samples during hyperparameter grid search. Models were evaluated across four core performance metrics defined on True Positives (TP), False Positives (FP), True Negatives (TN), and False Negatives (FN):

- **Accuracy** = $(\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN})$: Measures total overall classification correctness across safe and phishing instances.
- **Precision** = $\text{TP} / (\text{TP} + \text{FP})$: Measures the exact ratio of correctly flagged phishing emails relative to all emails flagged as phishing, quantifying false alarm risks.
- **Recall (Sensitivity)** = $\text{TP} / (\text{TP} + \text{FN})$: Measures the proportion of actual phishing emails correctly caught by the model, quantifying missed threat risks.
- **F1-Score** = $2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$: Computes the harmonic mean of Precision and Recall, serving as the primary benchmark metric.

TABLE III. HOLDOUT TEST SET PERFORMANCE (3,507 SAMPLES)

Model Architecture	Acc.	Prec.	Rec.	F1
Logistic Reg.	98.26%	97.28%	98.09%	97.68%
Random Forest	98.00%	97.04%	97.64%	97.34%
Naive Bayes	97.15%	96.54%	95.80%	96.17%
MLP Neural Net	98.29%	98.30%	97.10%	97.70%
Soft-Voting Ens.	98.52%	98.09%	97.94%	98.02%

Comparative Performance and Error Analysis

As detailed in Table III, the Soft-Voting Ensemble obtained the highest overall accuracy of 98.52% and an F1-Score of 98.02% on the independent holdout test set. Among individual base classifiers, the MLP Neural Network achieved 98.29% accuracy, leveraging dense non-linear feature representations in its 100-neuron hidden layer. Logistic Regression achieved 98.26% accuracy, demonstrating that L2-regularized linear models perform exceptionally well when supplied with scaled Word2Vec and metadata features.

Confusion Matrix Diagnostics

Examining confusion matrix counts on the 3,507 holdout test samples provides granular insight into misclassification behavior. Out of 3,507 samples (comprising 2,196 legitimate emails and 1,311 phishing emails), the Soft-Voting Ensemble misclassified only 25 legitimate emails as phishing (False Positives = 25) and missed only 27 phishing emails (False Negatives = 27), yielding a True Positive count of 1,284 and True Negative count of 2,171.




FIGURE 3. Confusion matrix for Soft-Voting Ensemble.




FIGURE 4. Confusion matrix for Random Forest classifier.




FIGURE 5. Confusion matrix for Logistic Regression model.




FIGURE 6. Confusion matrix for MLP Neural Network.




FIGURE 7. Confusion matrix for Multinomial Naive Bayes.

ROC Curves and Feature Drivers

Receiver Operating Characteristic (ROC) curve analysis evaluates model discrimination capability across all classification thresholds. As shown in Figure 8, the Soft-Voting Ensemble achieves an Area Under Curve (AUC) of 0.998, indicating near-perfect separability between safe and phishing instances. Figures 9 and 10 illustrate feature importance drivers, demonstrating that structural metadata indicators (such as url_count, urgency_score, caps_ratio, and punctuation_density) provide high-confidence signals alongside top TF-IDF n-grams.




FIGURE 8. Comparative ROC curves (Ensemble AUC = 0.998).




FIGURE 9. Random Forest top feature importance drivers.



FIGURE 10. Logistic Regression top feature coefficients.

GATEWAY THROUGHPUT & OPERATIONAL DISCUSSION

Experimental Benchmarking & Profiling Methodology

In high-volume enterprise mail security infrastructure, classification models must execute within strict sub-millisecond service-level agreements (SLAs) to prevent mail transfer agent (MTA) queue accumulation. We conducted comprehensive latency benchmarking on an Intel Core i7 dual-core processor operating at 2.80 GHz with 16 GB RAM, running Ubuntu Linux 22.04 LTS without dedicated GPU hardware acceleration. Benchmark timings were recorded across 1,000 continuous inference iterations to capture realistic operational variance.

The end-to-end processing pipeline time for each incoming email is decomposed into two distinct computational stages: 1) *Feature Extraction Latency (1.85 ms)*: Comprising BeautifulSoup HTML parsing (0.40 ms), TF-IDF n-gram vectorization across 5,000 terms (0.80 ms), Word2Vec CBOW embedding matrix mean vector averaging (0.65 ms), and 12 structural metadata calculations; and 2) *Model Forward Inference Latency (0.25 ms)*: Evaluating pre-extracted feature vectors against trained model coefficients.




FIGURE 11. Per-email inference latency benchmark (ms).

Comparative Model Profiling & Hardware Economics

As illustrated in Figure 11, individual base classifiers exhibit sub-millisecond execution: Multinomial Naive Bayes completes prediction in 0.12 ms/email, Logistic Regression in 0.28 ms/email, and the MLP Neural Network in 0.85 ms/email. The complete Soft-Voting Ensemble requires 2.10 ms per message. In stark contrast, fine-tuned transformer architectures (such as BERT-base or RoBERTa-large) require 50 ms to 200 ms per forward pass on CPU hardware.

This order-of-magnitude efficiency differential translates directly into substantial infrastructure cost savings. A single worker thread running our Soft-Voting Ensemble processes up to 476 emails per second (28,560 emails per minute). Deploying an 8-thread containerized microservice node achieves a peak throughput exceeding 3,800 emails per second (over 228,000 emails per minute), easily absorbing enterprise morning email burst volumes without incurring the steep capital and operational expenditures associated with GPU cluster provision (\$1,500–\$4,000/month per node).

Concurrency, Microservice Architecture & Load Balancing

For seamless integration into existing enterprise mail gateways (e.g., Postfix, Microsoft Exchange, Sendmail), the detection pipeline is encapsulated as a stateless RESTful microservice packaged inside a Docker container. An NGINX reverse proxy distributes incoming inspection requests across worker processes via round-robin load balancing. Because feature extraction and model inference are stateless, the architecture scales horizontally across arbitrary Kubernetes pods, providing fault-tolerant auto-scaling during high-volume spam campaigns.

Resilience Against Adversarial Evasion & Vocabulary Drift

A critical operational challenge facing email security gateways is vocabulary drift—the rapid alteration of keywords by threat actors to evade static filters. Our hybrid feature design provides multi-layered resilience against adversarial evasion:

- **Semantic Word2Vec Generalization:** Custom continuous-bag-of-words (CBOW) embeddings map unseen zero-day synonyms into dense 100-dimensional vector spaces, retaining high cosine similarity to known phishing terms.
- **Structural Metadata Invariance:** Formatting features such as `url_count`, `caps_ratio`, `digit_ratio`, and `urgency_score` operate independently of textual language, acting as immutable behavioral anchors.

- **Sanitization Hardening:** BeautifulSoup HTML parsing strips invisible comments (e.g., `U<!-- tag -->RGENT`) and normalizes character encoding before vectorization.

SOC Quarantine Workflows & Triaging Thresholds

To minimize human operator intervention in Security Operations Centers (SOCs), classification probabilities output by the Soft-Voting Ensemble are mapped into three automated action tiers:

- 1) *Auto-Quarantine* ($P(y=1) \geq 0.85$): Message is immediately isolated into SOC quarantine; sender IP is logged to local firewall blocklists.
- 2) *Warning Banner Insertion* ($0.50 \leq P(y=1) < 0.85$): Message is delivered to recipient inbox with a prominent visual warning header advising against clicking links.
- 3) *Clean Delivery* ($P(y=1) < 0.50$): Message is delivered directly to inbox without latency delay.

CONCLUSION & FUTURE DIRECTIONS

Comprehensive Summary of System Accomplishments

This project established a robust, end-to-end machine learning framework for proactive phishing email classification. By constructing a 5,112-dimensional hybrid feature matrix that integrates sparse sublinear TF-IDF n-grams (5,000 features), dense continuous-bag-of-words Word2Vec embeddings (100 features), and 12 dense structural metadata metrics, the architecture captures syntactic vocabulary patterns, underlying semantic intent, and behavioral formatting indicators simultaneously.

Rigorous empirical evaluation across 17,535 preprocessed email samples and an independent 20% holdout test set (3,507 samples) confirmed that the Soft-Voting Ensemble achieves state-of-the-art performance: 98.52% Accuracy, 98.09% Precision, 97.94% Recall, and 98.02% F1-Score with an Area Under ROC Curve (AUC) of 0.998. The model misclassified only 25 legitimate emails out of 2,196 and missed only 27 phishing emails out of 1,311, establishing exceptional balance between high security and low false alarm rates.

Architectural Trade-offs & Gateway Economics

A central finding of this research is that classical ensemble models paired with multi-modal feature engineering provide accuracy parity with heavy transformer language models (such as BERT or RoBERTa) while delivering a 50x to 100x improvement in computational efficiency. While transformers require specialized GPU clusters to process high-volume mail streams, our Soft-Voting Ensemble executes in 2.10 ms per email on standard CPU hardware. This sub-millisecond latency enables a single CPU gateway worker thread to evaluate up to 476 emails per second, offering enterprise security operations centers a cost-effective, low-carbon infrastructure solution for real-time perimeter protection.

Strategic Future Research Directions

To further advance real-time email security, future research will pursue four key technical trajectories: 1) Lightweight Hybrid Fallback Architecture (routing borderline confidence scores to DistilBERT); 2) Header Telemetry Feature Fusion (incorporating SPF, DKIM, and DMARC verification flags); 3) Automated Active Learning Loops (online retraining on user-flagged threat samples); and 4) Explainable AI Dashboards (integrating SHAP values for transparent security operator triage).

ACKNOWLEDGEMENT

I would like to express my sincere gratitude to the **Indian Institute of Computing and Technology (IICT)** for providing me with the opportunity to undertake this internship and gain valuable practical experience in the fields of Artificial Intelligence, Machine Learning, Natural Language Processing, and Cybersecurity. The internship offered an excellent platform to strengthen my technical skills through hands-on projects, research-oriented learning, and real-world problem solving.

I am especially grateful to **Dr. Ashok Gopalakrishnan** for his exceptional guidance, mentorship, and unwavering support throughout the internship. His vast knowledge, practical teaching approach, and dedication to student learning made a significant impact on my understanding of machine learning concepts and their real-world applications. During the training program, he patiently addressed the questions and doubts of every student, ensuring that complex topics became easy to understand. His guidance during the implementation of machine learning projects laid a strong foundation for my learning and greatly enhanced my confidence in applying machine learning techniques. Furthermore, his valuable suggestions and practical insights

helped me identify appropriate datasets from Kaggle, enabling me to successfully complete both internship projects with confidence and accuracy.

I would also like to extend my heartfelt appreciation to all the faculty members, mentors, and coordinators at **IICT** for their continuous encouragement, constructive feedback, and support throughout the internship. Their collective efforts created an enriching learning environment that encouraged curiosity, innovation, and independent problem-solving.

Finally, I express my sincere gratitude to **Jecrc University**, for providing me with a strong academic foundation and continuous encouragement to pursue practical learning opportunities. The knowledge, experience, and confidence gained during this internship have significantly contributed to my academic and professional development and will serve as a strong foundation for my future career in Artificial Intelligence, Machine Learning, and Cybersecurity.

REFERENCES

- [1] M. Sahami, S. Dumais, D. Heckerman, and E. Horvitz, "A Bayesian approach to filtering junk e-mail," in *AAAI Workshop on Learning for Text Categorization*, 1998, pp. 98–105.
- [2] G. Egozi and R. Verma, "Phishing Email Detection Using Robust NLP Techniques," in *Proc. IEEE Int. Conf. Data Mining Workshops (ICDMW)*, Singapore, 2018, pp. 7–12.
- [3] A. Al-Subaee and A. Al-Zahrani, "An NLP-Based Phishing Email Detection Model Using Machine Learning," in *Proc. 3rd IEEE Int. Conf. Comput. Appl. Inf. Security (ICCAIS)*, Riyadh, 2020, pp. 1–6.
- [4] S. Gedam and T. Shende, "Phishing/Spam Email Detection with Natural Language Processing and Machine Learning," *Int. Res. J. Adv. Eng. Sci.*, vol. 8, no. 2, pp. 145–151, 2023.
- [5] S. A. Salloum, T. Gaber, S. Vadera, and K. Shaalan, "Phishing Email Detection Using Natural Language Processing Techniques: A Literature Survey," *Procedia Comput. Sci.*, vol. 189, pp. 100–108, 2021.
- [6] P. Bountakas and C. Xenakis, "HELPHISH: Heuristics and Machine Learning for Phishing Email Detection," *Comput. Secur.*, vol. 110, p. 102435, Nov. 2021.
- [7] R. Hari Krishnan, A. R. Nair, V. Asokan, and P. R. Hari, "A Machine Learning Approach to Email Classification for Phishing Detection Using NLP," in *Proc. IEEE ICACCI*, 2023, pp. 1–8.
- [8] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Efficient Estimation of Word Representations in Vector Space," *arXiv preprint arXiv:1301.3781*, 2013.
- [9] J. Devlin, M. W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," in *Proc. NAACL-HLT*, 2019, pp. 4171–4186.
- [10] L. Breiman, "Random Forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [11] F. Pedregosa et al., "Scikit-learn: Machine Learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [12] C. D. Manning, P. Raghavan, and H. Schütze, *Introduction to Information Retrieval*. Cambridge University Press, 2008.
- [13] Y. Bengio, R. Ducharme, P. Vincent, and C. Jauvin, "A Neural Probabilistic Language Model," *Journal of Machine Learning Research*, vol. 3, pp. 1137–1155, 2003.
- [14] A. K. Jain, R. P. W. Duin, and J. Mao, "Statistical pattern recognition: A review," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 22, no. 1, pp. 4–37, 2000.
- [15] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning representations by back-propagating errors," *Nature*, vol. 323, no. 6088, pp. 533–536, 1986.